

# Towards a Facial Recognition Workstation

With Accompanying Examples

Matthew Shane Kingon

~~g12k1298~~

Submitted in partial fulfilment of the  
requirements for the degree of BACHELOR OF  
SCIENCE (HONOURS) [of Rhodes University]



is this  
wording  
correct?

**RHODES UNIVERSITY**  
*Where leaders learn*

Computer Science Department

Rhodes University

Grahamstown, South Africa

Supervisor: Prof EP Wentworth

## Abstract

Abstracty things

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

Abstracty abstracty abstracty

# ACM Classification

classificationy things

## Acknowledgements

I would like to acknowledge the financial and technical support of Telkom SA, Tellabs, Genband, Easttel, Bright Ideas 39, THRIP and NRF SA (UID 75107) through the Telkom Centre of Excellence in the Department of Computer Science at Rhodes University.

↑  
did you get support from them ?

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Illumination</b>	<b>3</b>
2.1	The Issue of Illumination . . . . .	3
2.2	Comparing Two Methods . . . . .	4
2.2.1	Histogram Equalization . . . . .	4
2.2.2	Summary of Histogram Equalization . . . . .	5
2.2.3	Mean Illumination Estimation . . . . .	6
2.2.4	Summary of Mean Illumination Estimation . . . . .	8
2.3	Experiment 01: Setup . . . . .	9
2.3.1	Experiment 01: Results . . . . .	11
2.3.2	Experiment 01: Discussion . . . . .	14
2.4	Experiment 02: Setup . . . . .	15
2.4.1	Experiment 02: Results . . . . .	15
2.4.2	Experiment 02: Discussion . . . . .	18
2.5	summary . . . . .	18
<b>3</b>	<b>Implementation</b>	<b>21</b>
<b>4</b>	<b>Results</b>	<b>23</b>
<b>5</b>	<b>Discussion</b>	<b>25</b>
<b>6</b>	<b>Conclusion</b>	<b>27</b>
<b>A</b>	<b>Appendages</b>	<b>29</b>
<b>References</b>		<b>29</b>



# Chapter 1

## Introduction



# Chapter 2

## Illumination

### 2.1 The Issue of Illumination

Lighting plays a very big role in all facial recognition, identification, or just about any image processing problem. ~~As even though~~, to the human eye, it ~~is~~ is almost undetectable irregularity in the world we observe. ~~While~~, to a computer, that must ~~look~~ inspect intently at each and every pixel within an image to determine what it is observing, even the slightest change in lighting makes ~~each~~ each pixel value change dramatically. This problem makes it a painful process to perform any image processing techniques.

can

To better appreciate the problem it should be noted that there are an abundance of illusory images that manage to fool the human brain. These examples are often very specifically designed for this very purpose. However, it does illustrate the point. The example shown in 2.1 may be an old one, but will still often fool most people you show it to; the two blocks with the orange dot are in fact, the same shade of grey. Some may see it right away but even so it takes effort to see it. This is due to the fact that your brain logically assumes that the lower dot is from the brighter block set and hence should be brighter than the darker block set around it (which it is. However, it is not brighter than all dark blocks in the image), this occurs because your brain doesn't take into account the shadow cast upon the blocks.

remove  
parentheses

This is fun, but I don't really  
see its relevance.

Don't discard it, though.

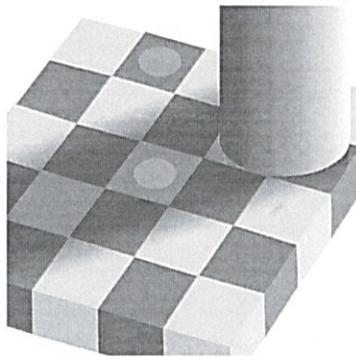


Figure 2.1: Optical Illusion example

When it comes to faces there are two main ways illumination could degrade our images. One, uniform change in ambient light this occurs when the light source is straight on, visually the whole image gets darker or brighter. Two, localised lighting changes, e.g. the light source is to the right of the face, so the left side is darker than the right. This is a distinction we will come back to later. Illumination is clearly something any face recognition service will have to take into account if they wish to accurately recognize a set of faces. Thus, illumination is a logical and readily available normalization target to test the ease of use of my system.

## 2.2 Comparing Two Methods

The plan will be to test a newly proposed method of illumination correction, The Mean Illumination Estimation (MIE) [1], against something with proven results, namely, OpenCV's built-in histogram equalization. We have the proposition put forward by [1] that their ~~a~~ Algorithm bests the current standard for illumination correction. ~~Thus implying it should fair better than the simple quick and easy OpenCV solution.~~ This implies that fair

### 2.2.1 Histogram Equalization

As stated we will be employing OpenCV's implementation of this ~~Algorithm~~ Algorithm. However, presented below is an overview on how the algorithm works and why it is theoretically useful to combat illumination conditions.

Correct for different

Before we get into the math of the method an intuitive description would be useful. ~~as such~~ What we are doing is taking the image and stretching the darkest pixel to 0 (pure black)

Since you're using latex,  
Formatting doing the maths decently  
should be easy.

and the brightest pixel to 255 (pure white) and the rest of the pixels distributed out in a similar manner in order to improve the overall contrast of the image. However, we don't affect the base structure of the image in doing so i.e. pixel  $(x_0, y_0)$  which is brighter than pixel  $(x_1, y_1)$  put through this algorithm to get  $(nx_0, ny_0)$  will still be brighter than  $(nx_1, ny_1)$  though both values will likely have changed.

Mathematically we are mapping a clustered distribution to a wider (0-255) more uniform spread of intensity values. Let  $f$  be the given image represented by a  $n \times m$  matrix of pixel intensities ranging from (0 to L-1) where 'L' is the No. of possible values (usually 256) Now 'p' will denote the normalized histogram of our image 'f' as a list of bins per intensity value:

$$p_n = \frac{x_n}{Total} \quad n = 0, 1, 2, \dots, L-2, L-1 \quad (2.1)$$

where  $x_n$  is the No. of pixels with intensity 'n' and 'Total' is the total No. of pixels in the image. We will define the Histogram equalized image as 'g';

$$g_{ij} = \text{floor}((L-1) \sum_{n=0}^{L-1} f_{ij} p_n) \quad (2.2)$$

Note; floor rounds down the value.

### 2.2.2 Summary of Histogram Equalization

Now we have a normalised image with which we can hopefully perform recognition on with greater success. We note that technically it doesn't really care that there is illumination issues present. Histogram equalization is a global image editor that will effect the entire image in the same way, regardless of the actual lighting conditions present in portions of the image. I.e. if we took a face image with one half bright and one half dark and put it through the algorithm, as already stated above, the brighter pixels on the one half will still be brighter after the equalization is performed which still equates to an illumination problem for the eigenface algorithm. However, these are very much edge cases for the most part illumination issues will be of the form of global lighting differences.

The key idea now is to estimate the regional value ~~effect~~ of the illumination, and ~~adjust for it~~. Use that ~~allow~~ to cancel out the illumination.

### 2.2.3 Mean Illumination Estimation

a localized

This method takes the smoothing approach to lighting normalization. This is done so as to remove the layer of the image responsible for illumination changes. Firstly it notes that

components according to the Illumination-Reflection model described in detail in [2], a pixel ~~in a facial~~ ~~image~~ ~~f<sub>xy</sub>~~ gets its value from two parts.  $r_{xy}$  represents the reflection component of an image at the point (x,y) and  $i_{xy}$  represents the illumination component. Thus we get the equation:

in a facial image

Components

$$f_{xy} = r_{xy} \times i_{xy} \quad (2.3)$$

Now, as  $r(x,y)$  is dependant purely on the surface material in question and not affected by illumination it would be an intrinsic representation of the facial image. Suppose  $i(x,y)$  changes little in value within a small area while in the presence of a weak light source. Thus we wish to find an estimate of our image  $f(x,y)$  that will allow us to do this separation. To attain such an estimate we apply a logarithmic transformation to our image  $f(x,y)$  we call this new function  $g(x,y)$

$$\begin{aligned} g_{xy} &= \ln(f_{xy}) \\ &= \ln(r_{xy}) + \ln(i_{xy}) \end{aligned} \quad (2.4)$$

Now the mean estimate for  $g_{xy}$  is obtained as:

$g_{xy}$  becomes a

$$\begin{aligned} \hat{g}_{xy} &= \frac{1}{n^2} \sum_{(s,t) \in \omega_{nn}} g_{st} \\ &= \frac{1}{n^2} \sum_{(s,t) \in \omega_{nn}} \ln(r_{st}) + \frac{1}{n^2} \sum_{(s,t) \in \omega_{nn}} \ln(i_{st}) \end{aligned} \quad (2.5)$$

Note  $\omega_{nn}$  is the area around a given pixel (x,y) with (s,t) being the enumeration of these pixels and n is the width/height of said kernel around (x,y). The Quotient image is constructed from equations (2),(3) we do so to eliminate  $i_{xy}$  (or make its contribution to the image negligible):

28

What is this?

$$\begin{aligned} d_{xy} &= g_{xy} - \hat{g}_{xy} \\ &= \ln(r_{xy}) - \frac{1}{n^2} \sum_{(s,t) \in \omega_{nn}} \ln(r_{st}) + \sigma \end{aligned} \quad (2.6)$$

Where:  $\sigma = \ln(i_{xy}) - (\frac{1}{n^2}) \sum_{(s,t) \in \omega_{nn}} \ln(i_{st})$  we note that  $\sigma$  will be a very small value and can hence be omitted from (4) leaving us with the relation:

$$\begin{aligned} d_{xy} &= g_{xy} - \hat{g}_{xy} \\ &\approx \ln\left(\frac{r_{xy}}{\left(\prod_{(s,t) \in \omega_{nn}} r_{st}\right)^{\frac{1}{n^2}}}\right) \end{aligned} \quad (2.7)$$

Now,  $d_{xy}$  represents the ratio between the current points reflectance and the average reflectance around it. when the materials are the same  $d_{xy}$  tends towards zero, but when they are different e.g. facial skin and facial features  $d_{xy}$  becomes notably non-zero. Let us consider:

$$\alpha = \frac{1}{ab} \sum_{(x,y) \in f_{a \times b}} |d_{xy}| \quad (2.8)$$

with  $a = \text{No. of rows}$  and  $b = \text{No. of columns}$  in an image  $f_{a \times b}$ .  $\alpha$  will represent the average grey value ratio of the facial skin and features the global difference is hence reduced as:

$$h_{xy} = \exp \frac{d_{xy}}{\alpha \beta} \quad (2.9)$$

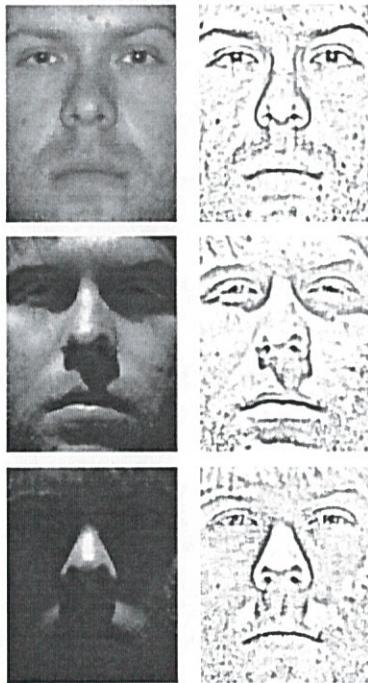
with  $\beta$  being a controllable scaling factor. However, it is usually set in the range of 2-3 thus, finally to highlight facial features and reduce impact of background noise, post processing is done as:

$$\hat{o}_{xy} = \begin{cases} h_{xy} & h_{xy} < 1 \\ 1 & h_{xy} \geq 1 \end{cases} \quad (2.10)$$

$$o_{xy} = \left[ \frac{255 \times (\hat{o}_{xy} - c)}{1 - c} \right] \quad (2.11)$$

With  $o_{xy}$  being the final result image to be used in training or recognition.  $c$  is the

minimum value of  $\hat{o}_{xy}$ . The math may be complex but the idea is rather simple, given an image, we separate out the intensity factor from the structure of the face, setting it to zero and rebuilding the face. This leaves us with an image that for the most part is devoid of all extra light. Ideally, multiple images of the same object under different lighting conditions that are put through this algorithm will end up looking the same. An example of it in use can be seen below (Note how, despite the drastic differences in illumination on the left side, the right side varies little from face to face.):



How big was  
n for this experiment  
and what size?  
are these images?

Figure 2.2: Mean Illumination Estimation

#### 2.2.4 Summary of Mean Illumination Estimation

The complications arise from the converting from a simple image  $f_{xy}$  into two separate images  $r_{xy}$  namely the facial structure/texture map and  $i_{xy}$  the intensity with the goal of removing  $i_{xy}$  from the equation. We are forced to achieve this by a logarithmic transform which only slightly (computationally) affects the image, the fact remains that it still does change the image.

You don't actually do the separation.  
Would it not be better to say "The complications  
arise from estimating the ~~the~~ illumination component so as  
to subtract it to highlight the remaining  
reflectance ..."

*we postulate that*

To start, we compare the output my MIE implementation provides against that of the output attained according to [1]. We note above in fig:[ 2.3];

We observe that my implementation provides near identical results to that of the article for the base face, but as the images become darker my ~~implementation seems to degrade~~. However, the artefacts are ~~too~~ uniform to be a natural error on the part of my implementation. Hence, I believe these artefacts are being generated from the conversion between image to PDF, and back to image again i.e. slight smudging is introduced in the dark regions from the conversion algorithm, possibly from compression algorithms.

*our results seem to deviate more from the published paper.*  
To Illustrate why I believe it is not my algorithm in the wrong here are several more images put through my algorithm that are not extracted from a PDF. *paper of [ref].*

~~None show~~ ~~similar~~  
~~artefacts~~

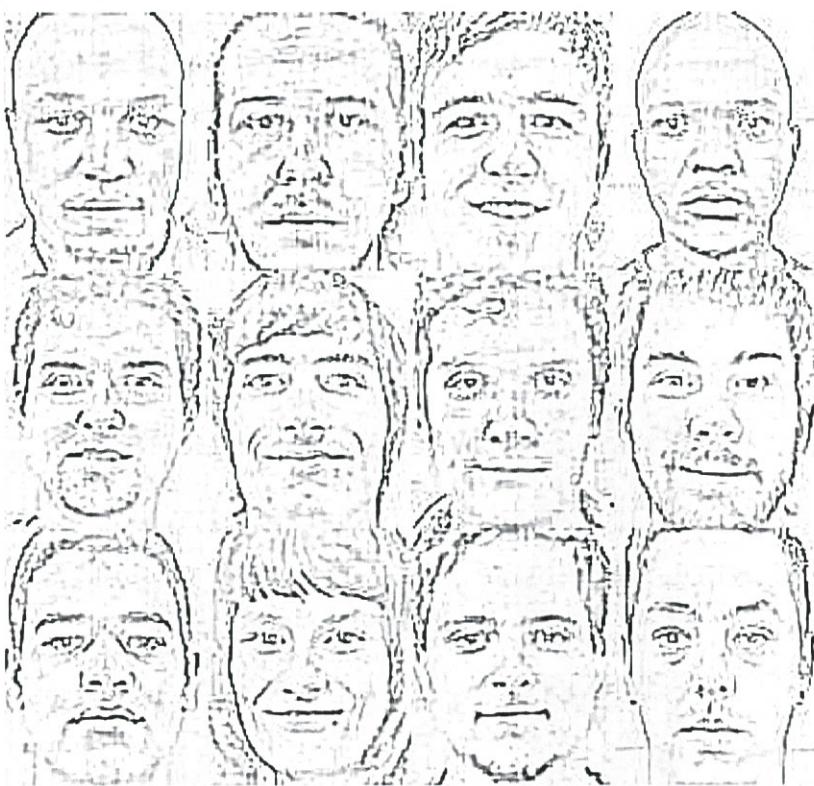


Figure 2.4: MIE implementation applied to CSC Honours 2015 Class list, Rhodes University

*To further investigate, I took the original PDF images and*  
Below fig:[ 2.5], we see the same faces as The MIE comparison test fig:[ 2.3] above. However, instead of the MIE applied to them they have undergone histogram equalization.

*ran them through the OpenCV  
histogram equalization . --- etc*

New paragraph.

Using the workstation, the task now is to gain more insight into the algorithm by attempting to duplicate the literature results, and to further understand the weaknesses and strengths

In conclusion, any attempt to normalize an image's illumination will undoubtedly degrade some aspects of the image we would rather retain. However, one would hope that the gained standardization of facial images out way this degradation and yield more accurate recognition rates. We also note that this method is a localised correctional algorithm. That is to say the afore mentioned issue with histogram equalization not being very effective for images with half bright, half dark faces doesn't apply here. The formula and reasoning were put forward by and learned from [1] an article that attempts to find a better way of solving the lighting issue with positive results for their effort. results i will be attempting to duplicate and backup, hopefully providing more insight into the algorithm's strengths and weaknesses.

## 2.3 Experiment 01: Setup

Before the experiment can be run we first attempt to confirm as many aspects of the system work as intended. This is done to provide validation to the results obtained from the experimental procedure. Though we predominantly wish to do so for our implementation of the Mean Illumination Estimation, for completion, we also confirm OpenCv's built in Histogram Equalization method works as intended.

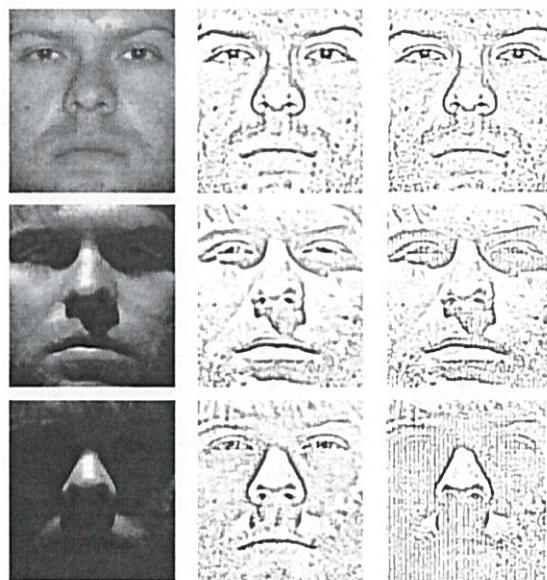


Figure 2.3: My MIE implementation [col 3] tested against provided example from [1] [col 2]



Figure 2.5: Histogram Equalization applied to article example faces

We see that the intensities have been stretched for all images, improving visibility. We also note the far right image, once again the same artefact pattern emerges. Pointing again to an issue with the image itself and not the algorithms performed upon them. Now that we have confidence in our implemented code we can construct an experiment to test the effects on accuracy they have.

To that end, experiment conditions will be; Yale B face database which contains 38 individuals with 65 images per subject (2470 total). However, to reduce computation time I will only use a subset of the data, the first 20 subjects. Note, the eigenface algorithm cannot handle images of varying size, so many of the ambient images need to be discarded as many are larger than the rest of the images. Also, some of the images were provided as corrupt, presumably as further testing data, they too have been removed. Furthermore we also take 10 images per subject to use as the training set. What remains is 20 subjects and 1072 images as testing data. The results we will be comparing are the total No of correct connections made by the eigenface algorithm. As noted above we expect MIE to surpass OpenCV's Histogram equalizations recognition accuracy.

### 2.3.1 Experiment 01: Results

We start with a base line, no normalization techniques used, just the plain dataset run through the eigenface algorithm. Doing so we get the graph;

Give more context here.

The algorithm "classifies" a face provided the writer gets a "deviation" score less than the threshold.

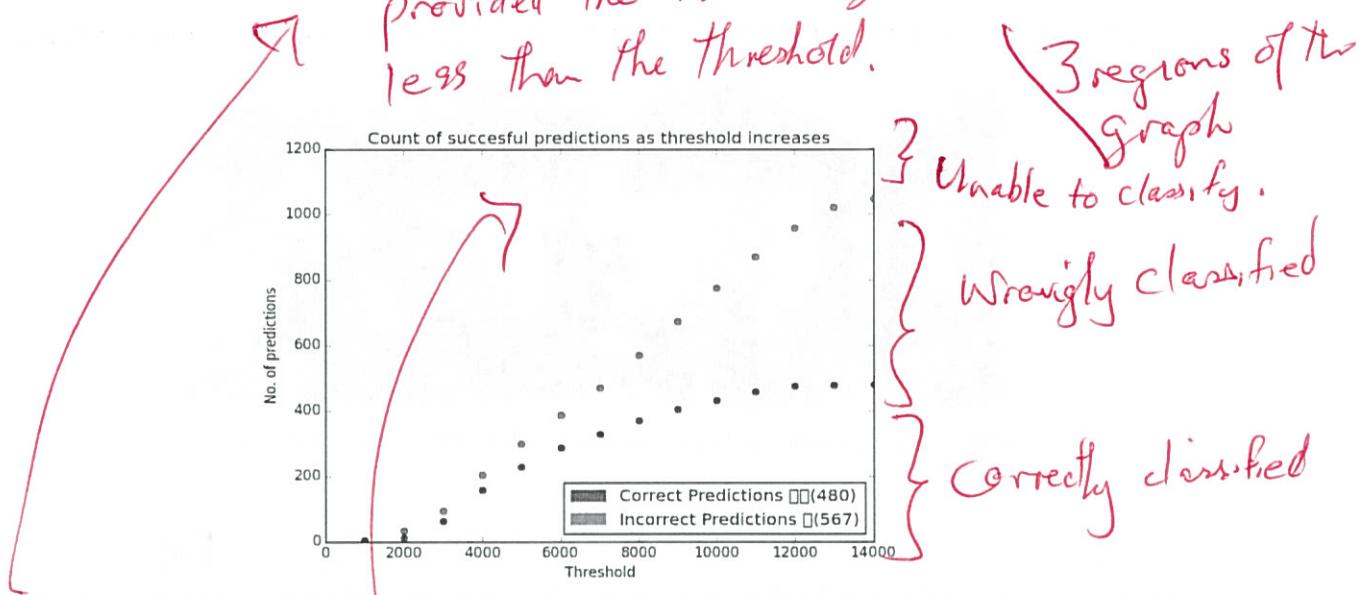


Figure 2.6: Base line, no normalization used

From Figure: 2.6 we see that we correctly recognize  $\frac{480}{1072}$  subject images i.e. roughly 45% accuracy. We run the experiment again but this time perform histogram equalization upon both test and training data. This provides us with the graph;

Can we force  
this grapher/  
edit the  
graph to  
draw to line  
at 1072

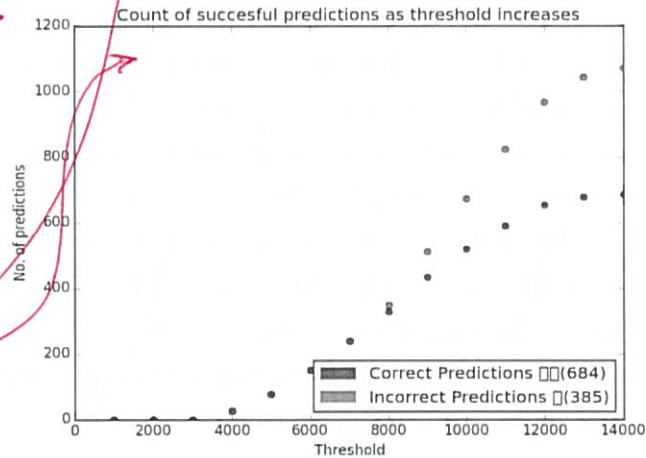


Figure 2.7: Histogram experiment

We now see from Figure: 2.7 that the system can correctly recognize  $\frac{684}{1072}$  of the subject images, roughly 64% accuracy rating, a meaningful improvement on the base line.

We now run the experiment again. However, instead of Histogram Equalization, we test our MIE algorithm. Doing so, we obtain the results;

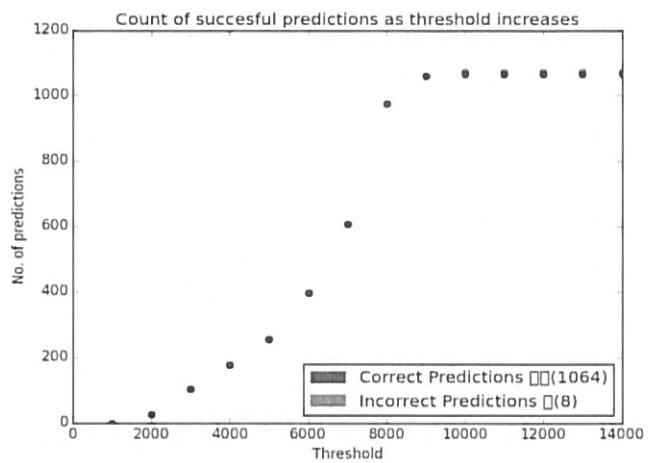


Figure 2.8: MIE experiment

Here we see greatly improved results with,  $\frac{1064}{1072}$  recognized faces, equating to roughly 99% accuracy.

In a final attempt to yield even better results, we ~~try~~ two more experiments, 1.) we run histogram equalization first, then MIE upon this normalised data. 2.) the other way round, ~~test~~ MIE first, then histogram equalization;

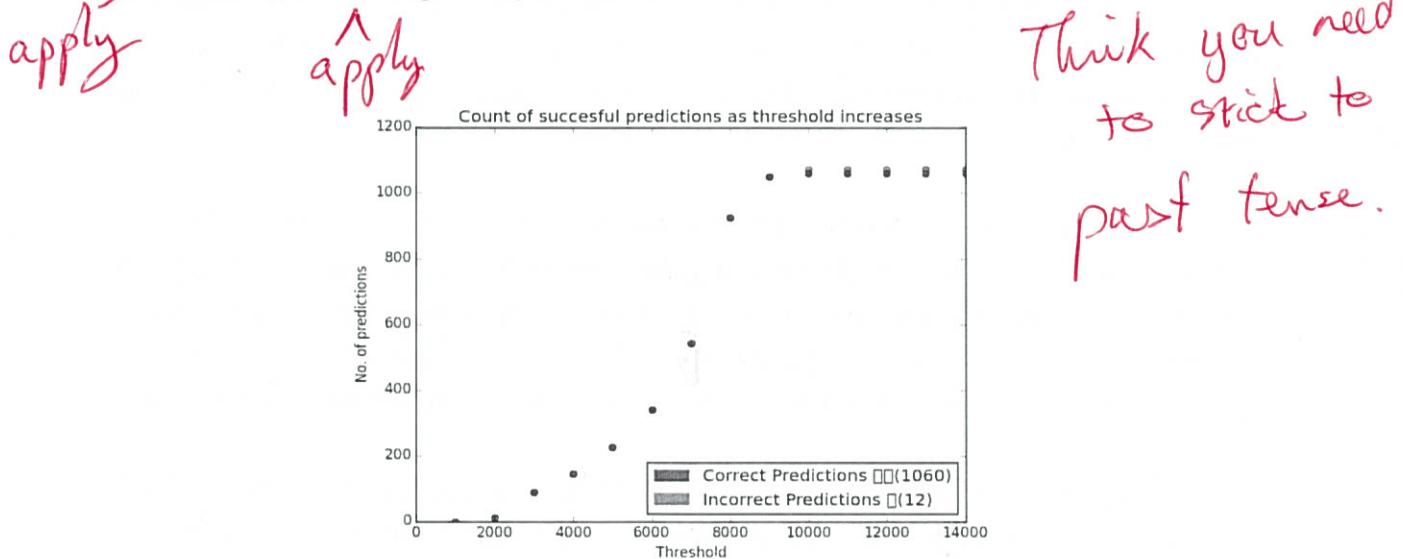


Figure 2.9: Histogram then MIE experiment

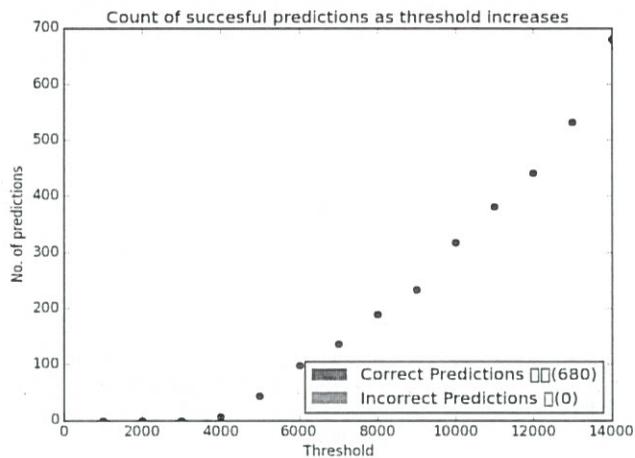


Figure 2.10: MIE then Histogram experiment

We see from these experiments that Performing Histogram Equalization first and then MIE on top of that performs nearly as well as MIE on its own, but the other way round yields shocking results, failing to even score all images under the maximum 14000 threshold.

→ No graph?

### 2.3.2 Experiment 01: Discussion

The Baseline experiment, fig:[ 2.6] provides a very poor score with this facial database, this was to be expected due to the high impact of lighting conditions on the eigenface algorithm, yielding a score of  $\frac{480}{1072}$ , roughly 45% accuracy. On top of that, it technically doesn't even score all the images, 3 where missed I.e. their threshold values where higher than the upper limit of our graph.

The first experiment we try is histogram equalization. Here we see an improvement in recognition at roughly 64%. We already know from the explanation of the workings of histogram equalization that it is not ideal for the eigenface method when used on images with localised lighting conditions something that the Yale B database has in abundance. This aside, all in all, it is a notable improvement on the none-normalised score but still far from the ideal.

Next we hope to achieve better results with the MIE version of the data. We get results which far surpass those provided by histogram equalization, at roughly 99%. This is because the MIE is a very Localised oriented Algorithm and as the Yale B database has nearly insignificant change in orientation and angle, purely lighting changes, the MIE performs

*well*

~~phenomenally~~ As such, these results most certainly agree with the claim published in [1].

With regards to the compound tests I deduce the reason for such discrepancy between the order is that performing histogram equalization first outputs a very lightly affected image with just a simple contrast change that MIE can operate on in much the same way as before, while trying to do it the other way, MIE outputs a vastly different image from the original, more resembling a sketch than a photo, with large variation in colour intensities, as such the histogram equalization algorithm falters as we have reduced the colour space to an even more compact region of the spectrum. Thus when we pass these new images through the histogram equalization method it just does more damage by stretching it further.

In reality we probably wouldn't want to perform multiple normalization techniques that target the same issue upon our dataset as we would expect such conflicts to arise. On the other hand we would hope that combining these techniques with others that target different aspects of normalization would not conflict with each other. Indeed, we would hope that doing so would provide a higher degree of accuracy in recognition.

## 2.4 Experiment 02: Setup

Next we will test our system against the AT&T database, this database contains 40 subjects with 10 images per subject totalling 400 images. We have evenly split this data into training and test data I.e. 5 images per subject goes to training and the other five to test data. Once again, the results we will be comparing are the total No. of correct connections made by the eigenface algorithm. This database primarily focuses on orientation of the subjects faces and has limited Illumination variance as such we expect the eigenface algorithm to perform with some degree of accuracy and the application of our two normalisation techniques to have minimal if any effect on the accuracies obtained.

### 2.4.1 Experiment 02: Results

As before we start with a base line, no normalization techniques used, just the plain dataset run through the eigenface algorithm. Doing so we get the graph;

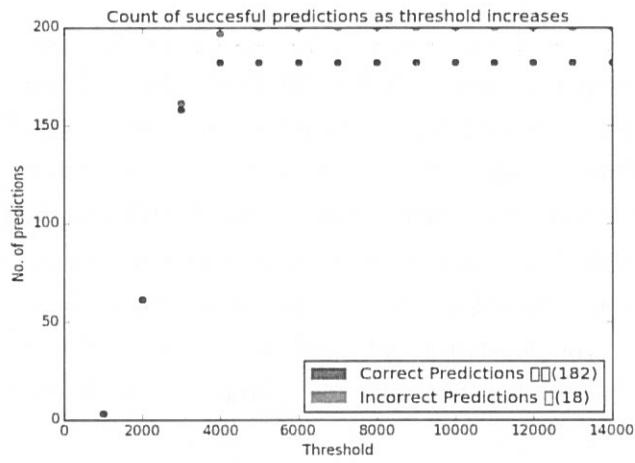


Figure 2.11: Base line, no normalization used

From Figure: 2.11 we see that we correctly recognize  $\frac{182}{200}$  subject images i.e. equating to a 91% accuracy. We run the experiment again but this time perform histogram equalization upon both test and training data. This provides us with the graph;

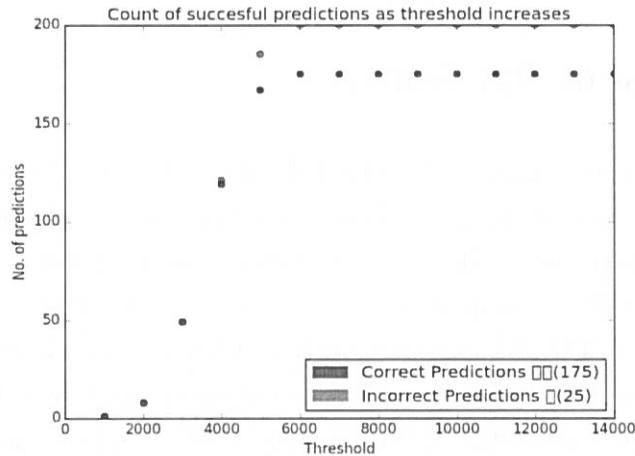


Figure 2.12: Histogram experiment

We now see from Figure: 2.12 that the system can correctly recognize  $\frac{175}{200}$  of the subject images, roughly 87% accuracy rating, a slight drop from the base line but still a viable accuracy.

We now run the experiment again. However, now we test our MIE algorithm again. Doing so, we obtain the results;

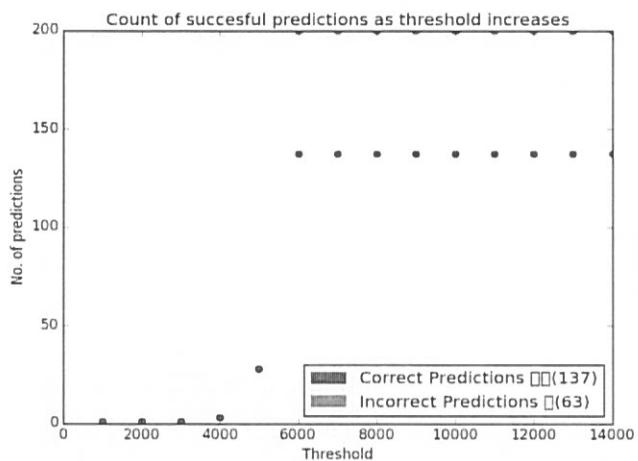


Figure 2.13: MIE experiment

Here we see greatly diminished results with,  $\frac{137}{200}$  recognized faces, equating to roughly 68% accuracy.

Once again, we try two more experiments, 1.) we run histogram equalization first, then MIE upon this normalised data. 2.) the other way round, test MIE first, then histogram equalization;

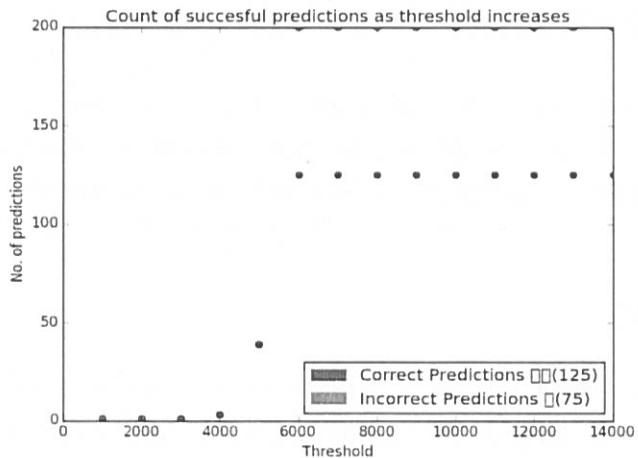


Figure 2.14: Histogram then MIE experiment

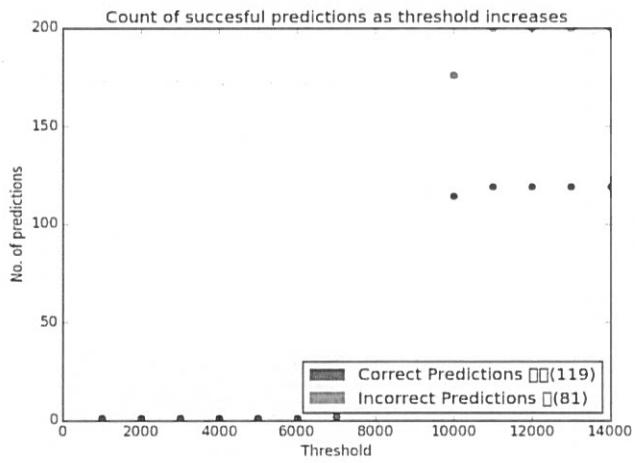


Figure 2.15: MIE then Histogram experiment

Again, aside from the over all inaccuracy of the combination we see that the previous pattern emerges hist then MIE performs in accordance with MIE on its own but the other way round provides even worse results.

#### 2.4.2 Experiment 02: Discussion

Initially we observed what we expected, the baseline performed with a high degree of accuracy due to the lack of illumination change in the images. However, to our disappointment when we run our correction algorithms we see diminished results with MIE being completely impractical for normal use.

What we take away from this is that although MIE performs phenomenally against strong illumination conditions it tends to degrade the image it is working on to an extent that other normalization techniques are ineffective at best and further degrading at worst.

### 2.5 Summary

So, once again we can clearly see that when it comes to illumination the MIE out performs OpenCV's histogram equalization by a large margin. Indeed, I find it hard to believe any algorithm could improve upon those accuracies. Sadly that's where the downsides must be mentioned, the algorithm ~~doesn't~~ play well with other techniques and even less so when there

are no illumination issues present, tending to be more of a hindrance than a help. Also, while ~~Accurate~~ the Python implementation was horribly time inefficient. As a benchmark, OpenCV's Histogram equalisation was performed upon the 1272 and did so in under a minute while ~~My~~ Python Implementation of MIE took roughly 10hrs to do the same ~~No.~~ of images. ~~the~~ ~~my~~ ~~A~~ But again, this speed difference comes with the benefit of being a very localised method. Each pixel is compared to a 11x11 kernel around itself as well as the average intensity for the whole image. Meaning it can handle images which are dark on one side but light on the other without any ~~hassle~~. ~~This means~~ problems. This means.

This per pixel kernel methodology indicates that this method could see a large speed-up if performed through a GPU. However, that is an idea for future work to look into.

~~was beyond the scope of this work to attempt to optimize this method. However, we do note~~

the following : ① Coding the existing logic into C++ or C is possible in Python, and should offer substantial speed-ups

② The current ~~Our~~ implementation is fairly naive: the 11 x 11 region is recomputed for every pixel in the image. Many efficient image processing algorithms operate by sliding a difference window over the image, and computing a running estimate of local intensity. There are undoubtedly gains performance gains to be had by doing this.

19 GPU.

PTO



## Chapter 3

### Implementation

I think your chapter would be better if you put the experiments in the opposite order.

"Bad results first", "Victory later"  
and then stronger take-away conclusions.  
"Yes, MIE looks good, but only for a narrow category of lighting problems."

Then contextualize: for classroom captured faces from a video stream, for our purposes, we don't anticipate highly directional or partially lit faces - the general ambient lighting in the venue is probably more significant, so we believe the histogram equalizer in OpenCV will be the algorithm of choice.



# Chapter 4

## Results



# Chapter 5

## Discussion



# Chapter 6

## Conclusion



## **Appendix A**

### **Appendages**



## Bibliography

- [1] Y.-P. G. Yong Luo and C.-Q. Zhang, "A robust illumination normalization method based on mean estimation for face recognition," *ISRN Machine Vision*, vol. 2013, no. 516052, pp. 0–10, 2013.
- [2] J. Ho, B. V. Funt, and M. S. Drew, "Separating a color signal into illumination and surface reflectance components: Theory and applications," *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, vol. 12, no. 10, pp. 966–977, 1990.

do they work with  
greyscale ?

